

# DistractedDriverDetection

May 16, 2025

```
[1]: # importing libraries
import os
import numpy as np
import pandas as pd
import tensorflow as tf
```

```
[2]: # loading dataset
data = pd.read_csv("/kaggle/input/state-farm-distracted-driver-detection/
↳driver_imgs_list.csv");
print(data.shape);
print(data.head());
```

(22424, 3)

	subject	classname	img
0	p002	c0	img_44733.jpg
1	p002	c0	img_72999.jpg
2	p002	c0	img_25094.jpg
3	p002	c0	img_69092.jpg
4	p002	c0	img_92629.jpg

```
[3]: # function to read image
def readimage(path):
    img = tf.keras.utils.load_img(path,color_mode = "rgb");
    img = tf.keras.utils.img_to_array(img);
    return img/255;
```

```
[4]: # plotting training images
from matplotlib import pyplot as plt
baseurl = "/kaggle/input/state-farm-distracted-driver-detection/imgs/train/";
activitymap = {'c0' : 'Safe driving','c1': 'Texting - right','c2': 'Talking on
↳the phone - right',
'c3': 'Texting - left','c4': 'Talking on the phone - left','c5': 'Operating the
↳radio',
'c6': 'Drinking','c7': 'Reaching behind','c8': 'Hair and makeup','c9': 'Talking
↳to passenger'};
fig = plt.figure(figsize = (10,4));
for i,directory in enumerate(sorted(activitymap.keys())):
    imagefile = os.listdir(os.path.join(baseurl,directory))[0];
```

```

imagepath = os.path.join(baseurl,directory,imagefile);
image = readimage(imagepath);
ax = fig.add_subplot(2,5,i + 1);
ax.imshow(image);
ax.set_title(activitymap[directory]);
ax.axis("off");
plt.tight_layout();

```



```

[5]: # importing libraries and setting up neural network model
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense,Conv2D,Flatten,MaxPooling2D,Dropout
model = Sequential();

```

```

[6]: # adding layers to neural network model
model.add(Conv2D(128,(3,3),activation = "relu",input_shape = (240,240,3),
data_format = "channels_last"));
model.add(MaxPooling2D((2,2)));
model.add(Conv2D(64,(3,3),activation = "relu"));
model.add(MaxPooling2D((2,2)));
model.add(Conv2D(32,(3,3),activation = "relu"));
model.add(MaxPooling2D((2,2)));
model.add(Flatten());
model.add(Dense(1024,activation = "relu"));
model.add(Dropout(0.2));
model.add(Dense(1024,activation = "relu"));
model.add(Dropout(0.2));
model.add(Dense(256,activation = "relu"));
model.add(Dropout(0.2));
model.add(Dense(10,activation = "softmax"));

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential

models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[7]: # compiling the model
model.compile(optimizer = "adam", loss = "categorical_crossentropy",
metrics = ["accuracy", tf.keras.metrics.Precision(), tf.keras.metrics.Recall(), tf.
↳keras.metrics.AUC()]);
```

```
[8]: # Import ImageDataGenerator and define paths, generators for training,
↳validation, and test sets
from tensorflow.keras.preprocessing.image import ImageDataGenerator
train_dir = "/kaggle/input/state-farm-distracted-driver-detection/imgs/train";
test_dir = "/kaggle/input/state-farm-distracted-driver-detection/imgs/";
train_datagen = ImageDataGenerator(rescale = 1.0/255,
shear_range = 0.2,
zoom_range = 0.2,
horizontal_flip = True,
validation_split = 0.2);
train_generator = train_datagen.flow_from_directory(
train_dir, target_size = (240, 240), class_mode = "categorical", batch_size =
↳32,
subset = "training"
);
validation_generator = train_datagen.flow_from_directory(
train_dir, target_size = (240, 240), class_mode = "categorical", batch_size =
↳32,
subset = "validation"
);
test_datagen = ImageDataGenerator(rescale = 1.0/255);
test_generator = test_datagen.flow_from_directory(
test_dir, target_size = (240, 240), classes = ["test"], class_mode =
↳None, batch_size = 32, shuffle = False
)
```

Found 17943 images belonging to 10 classes.

Found 4481 images belonging to 10 classes.

Found 79726 images belonging to 1 classes.

```
[9]: # using earlystopping
from tensorflow.keras.callbacks import EarlyStopping
earlystopping = EarlyStopping(monitor = "val_accuracy", patience = 5, min_delta =
↳0.005,
restore_best_weights = True);
```

```
[10]: # fitting the model
history = model.fit(train_generator, epochs = 10, validation_data =
↳validation_generator, callbacks = [earlystopping], verbose = 2)
```

Epoch 1/10

/usr/local/lib/python3.10/dist-

packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:122:

UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self.\_warn\_if\_super\_not\_called()

561/561 - 347s - 618ms/step - accuracy: 0.3872 - auc: 0.8175 - loss: 1.6514 - precision: 0.7688 - recall: 0.2294 - val\_accuracy: 0.7353 - val\_auc: 0.9687 - val\_loss: 0.7484 - val\_precision: 0.8307 - val\_recall: 0.6340

Epoch 2/10

561/561 - 299s - 534ms/step - accuracy: 0.8355 - auc: 0.9831 - loss: 0.5035 - precision: 0.8887 - recall: 0.7889 - val\_accuracy: 0.9105 - val\_auc: 0.9930 - val\_loss: 0.2839 - val\_precision: 0.9317 - val\_recall: 0.8924

Epoch 3/10

561/561 - 291s - 518ms/step - accuracy: 0.9180 - auc: 0.9933 - loss: 0.2695 - precision: 0.9362 - recall: 0.9008 - val\_accuracy: 0.9306 - val\_auc: 0.9956 - val\_loss: 0.2309 - val\_precision: 0.9446 - val\_recall: 0.9127

Epoch 4/10

561/561 - 291s - 519ms/step - accuracy: 0.9395 - auc: 0.9955 - loss: 0.2013 - precision: 0.9512 - recall: 0.9264 - val\_accuracy: 0.9493 - val\_auc: 0.9970 - val\_loss: 0.1639 - val\_precision: 0.9575 - val\_recall: 0.9409

Epoch 5/10

561/561 - 291s - 519ms/step - accuracy: 0.9510 - auc: 0.9964 - loss: 0.1666 - precision: 0.9613 - recall: 0.9434 - val\_accuracy: 0.9580 - val\_auc: 0.9967 - val\_loss: 0.1418 - val\_precision: 0.9660 - val\_recall: 0.9514

Epoch 6/10

561/561 - 292s - 520ms/step - accuracy: 0.9592 - auc: 0.9973 - loss: 0.1357 - precision: 0.9665 - recall: 0.9523 - val\_accuracy: 0.9621 - val\_auc: 0.9973 - val\_loss: 0.1298 - val\_precision: 0.9695 - val\_recall: 0.9574

Epoch 7/10

561/561 - 313s - 559ms/step - accuracy: 0.9637 - auc: 0.9974 - loss: 0.1211 - precision: 0.9702 - recall: 0.9580 - val\_accuracy: 0.9565 - val\_auc: 0.9975 - val\_loss: 0.1375 - val\_precision: 0.9655 - val\_recall: 0.9502

Epoch 8/10

561/561 - 298s - 532ms/step - accuracy: 0.9678 - auc: 0.9980 - loss: 0.1055 - precision: 0.9730 - recall: 0.9637 - val\_accuracy: 0.9587 - val\_auc: 0.9963 - val\_loss: 0.1397 - val\_precision: 0.9646 - val\_recall: 0.9547

Epoch 9/10

561/561 - 295s - 526ms/step - accuracy: 0.9702 - auc: 0.9980 - loss: 0.1007 - precision: 0.9755 - recall: 0.9663 - val\_accuracy: 0.9734 - val\_auc: 0.9982 - val\_loss: 0.0942 - val\_precision: 0.9775 - val\_recall: 0.9696

Epoch 10/10

561/561 - 291s - 518ms/step - accuracy: 0.9738 - auc: 0.9985 - loss: 0.0907 - precision: 0.9779 - recall: 0.9696 - val\_accuracy: 0.9748 - val\_auc: 0.9985 - val\_loss: 0.0795 - val\_precision: 0.9784 - val\_recall: 0.9726

```
[11]: # evaluating the model on validation data
print(model.evaluate(validation_generator));
```

```
141/141          58s 411ms/step -
accuracy: 0.9747 - auc: 0.9986 - loss: 0.0950 - precision: 0.9786 - recall:
0.9724
[0.08912034332752228, 0.9752287268638611, 0.9788906574249268,
0.9727739095687866, 0.9984093904495239]
```

```
[12]: # printing model summary
print(model.summary());
```

Model: "sequential"

Layer (type)	Output Shape	
Param #		
conv2d (Conv2D)	(None, 238, 238, 128)	
↳ 3,584		
max_pooling2d (MaxPooling2D)	(None, 119, 119, 128)	
↳ 0		
conv2d_1 (Conv2D)	(None, 117, 117, 64)	
↳ 73,792		
max_pooling2d_1 (MaxPooling2D)	(None, 58, 58, 64)	
↳ 0		
conv2d_2 (Conv2D)	(None, 56, 56, 32)	
↳ 18,464		
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 32)	
↳ 0		
flatten (Flatten)	(None, 25088)	
↳ 0		
dense (Dense)	(None, 1024)	
↳ 25,691,136		
dropout (Dropout)	(None, 1024)	
↳ 0		
dense_1 (Dense)	(None, 1024)	
↳ 1,049,600		

dropout_1 (Dropout)	(None, 1024)	⌵
↪ 0		
dense_2 (Dense)	(None, 256)	⌵
↪ 262,400		
dropout_2 (Dropout)	(None, 256)	⌵
↪ 0		
dense_3 (Dense)	(None, 10)	⌵
↪ 2,570		

Total params: 81,304,640 (310.15 MB)


Trainable params: 27,101,546 (103.38 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 54,203,094 (206.77 MB)

None

```
[14]: # plotting accuracy and loss curves
from matplotlib import pyplot as plt
fig,ax = plt.subplots(2,1,figsize = (18,10));
ax[0].plot(history.history["loss"],color = "b",label = "Training Loss");
ax[0].plot(history.history["val_loss"],color = "r",label = "Validation Loss");
ax[0].set_xlabel("Epochs");
ax[0].set_ylabel("Loss");
ax[0].set_title("Loss Curves");
ax[0].legend(loc = "best",shadow = True);
ax[1].plot(history.history["accuracy"],color = "b",label = "Training Accuracy");
ax[1].plot(history.history["val_accuracy"],color = "r",label = "Validation_
↪Accuracy");
ax[1].set_xlabel("Epochs");
ax[1].set_ylabel("Accuracy");
ax[1].set_title("Accuracy Curves");
ax[1].legend(loc = "best",shadow = True);
plt.tight_layout();
```



```
[26]: # predicting using test data
ypred = model.predict(test_generator);
print(ypred.shape);
```

2492/2492                      313s 125ms/step  
(79726, 10)

```
[29]: # processing for submission
filenames = test_generator.filenames
imageids = [f.split("/")[-1] for f in filenames];
```

```
[32]: # processing for submission
df = pd.DataFrame(ypred, columns = [f"c{i}" for i in range(10)])
df.insert(0, "img", imageids);
```

```
[33]: # Creating Results.csv
df.to_csv("Results.csv", index = False);
```

```
[40]: # generating validation data
validation_generator = train_datagen.flow_from_directory(
    train_dir, target_size = (240, 240), class_mode = "categorical", batch_size = 32,
    subset = "validation", shuffle = False
);
```

Found 4481 images belonging to 10 classes.

```
[41]: # changing ypred and ytest to normal form
ypred = model.predict(validation_generator);
ypred = np.argmax(ypred,axis = 1);
ytest = validation_generator.classes;
print(ypred.shape,ytest.shape);
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
    self._warn_if_super_not_called()


141/141          57s 407ms/step
(4481,) (4481,)
```

```
[44]: # printing accuracy score
from sklearn.metrics import
    accuracy_score,confusion_matrix,classification_report
score = accuracy_score(ytest,ypred);
print(score);
```

```
0.9750055791118054
```

```
[48]: # printing confusion matrix
import seaborn as sns
from matplotlib import pyplot as plt
arr = [f"c{i}" for i in range(10)];
plt.figure(figsize = (8,6));
cm = confusion_matrix(ytest,ypred);
sns.heatmap(cm,cmap = "Blues",fmt = "d",annot = True,xticklabels =
    arr,yticklabels = arr);
plt.xlabel("Predicted Values");
plt.ylabel("Actual Values");
plt.title("Confusion Matrix");
```





```
[49]: # printing classification report
report = classification_report(ytest,ypred,output_dict = True,target_names = _
    ↪arr);
df = pd.DataFrame(report).transpose();
plt.figure(figsize = (10,5));
sns.heatmap(df,cmap = "Blues",fmt = ".2f",annot = True);
plt.title("Classification Report");
```

